

notebook

April 28, 2025

```
[192]: import pandas as pd
import re
from sklearn.preprocessing import OrdinalEncoder, LabelEncoder
from typing import Tuple
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import time
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, \
    GradientBoostingClassifier, ExtraTreesClassifier
from lightgbm import LGBMClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, classification_report
```

```
[193]: data = pd.read_excel("./data/data.xlsx")
```

```
[194]: data.head()
```

```
[194]:
```

	explanation_quality	final_grade	study_hours	absences	group_work	schedule	\
0	5	9.46	4.3	<25%	oui	matin	
1	9	5.13	21.3	<25%	Non	matin	
2	7	9.67	18.9	<50%	Non	matin	
3	6	13.24	0.8	<25%	oui	matin	
4	3	12.28	48.0	<50%	oui	matin	

	class_participation	exam_content	\
0	un petit peu	mélange	
1	moyen	fait parti du cours enseigner par le prof	
2	un petit peu	fait parti du cours enseigner par le prof	
3	un petit peu	mélange	
4	un petit peu	fait parti du cours enseigner par le prof	

```

            result
0          non valide
1          non valide
2            valide
3  valide apres rattrapage
4            valide

```

```
[195]: data.describe()
```

```

[195]:      explanation_quality  final_grade  study_hours
count          5000.000000    5000.000000    5000.000000
mean             5.343800     10.965118     23.969980
std              2.079631      2.536466      8.422406
min              1.000000      5.000000      0.000000
25%              4.000000      9.240000     19.400000
50%              5.500000     10.960000     24.600000
75%              7.000000     12.700000     29.700000
max             10.000000     20.000000     50.000000

```

```
[196]: data.isnull().sum()
```

```

[196]: explanation_quality    0
final_grade                 0
study_hours                 0
absences                    0
group_work                  0
schedule                    0
class_participation         0
exam_content                0
result                      0
dtype: int64

```

1 EDA

```

[197]: # How many rows and columns
data.shape

```

```
[197]: (5000, 9)
```

```

[198]: # Quick look at the first few rows
data.head()

```

```

[198]:      explanation_quality  final_grade  study_hours  absences  group_work  schedule  \
0                5          9.46          4.3    <25%      oui    matin
1                9          5.13         21.3    <25%     Non    matin
2                7          9.67         18.9    <50%     Non    matin

```

3	6	13.24	0.8	<25%	oui	matin
4	3	12.28	48.0	<50%	oui	matin

	class_participation	exam_content	\
0	un petit peu	mélange	
1	moyen	fait parti du cours enseigner par le prof	
2	un petit peu	fait parti du cours enseigner par le prof	
3	un petit peu	mélange	
4	un petit peu	fait parti du cours enseigner par le prof	

	result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage
4	valide

```
[199]: # Detailed info about types and missing values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   explanation_quality    5000 non-null   int64
1   final_grade            5000 non-null   float64
2   study_hours            5000 non-null   float64
3   absences               5000 non-null   object
4   group_work             5000 non-null   object
5   schedule               5000 non-null   object
6   class_participation    5000 non-null   object
7   exam_content           5000 non-null   object
8   result                 5000 non-null   object
dtypes: float64(2), int64(1), object(6)
memory usage: 351.7+ KB
```

```
[200]: print(data['study_hours'].isna().sum())
```

0

```
[201]: # List of your categorical columns
categorical_cols = ['group_work', 'absences', 'schedule',
                    'class_participation', 'exam_content', 'result']

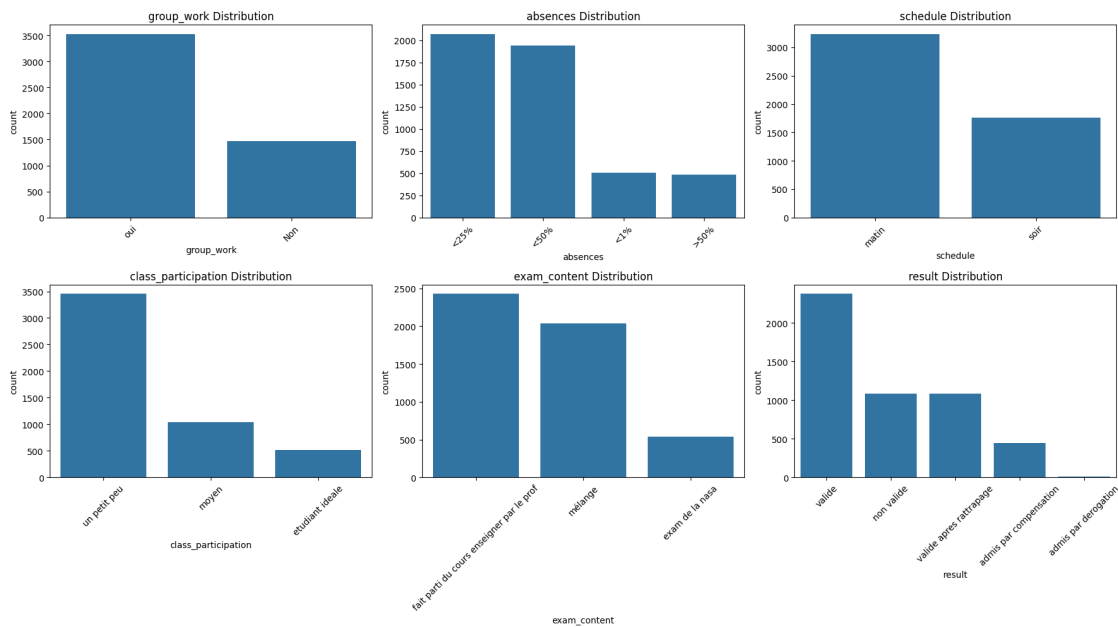
# Create subplots
plt.figure(figsize=(18, 10)) # Big figure for multiple plots
```

```

for idx, col in enumerate(categorical_cols):
    plt.subplot(2, 3, idx + 1) # 2 rows, 3 columns, plot number idx+1
    sns.countplot(x=col, data=data, order=data[col].value_counts().index)
    plt.title(f"{col} Distribution")
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

```



2 Interpretation of Categorical Feature Distributions

2.1 1. group_work Distribution

- “Non” (no group work) is more frequent than “oui” (yes group work).
- Students who did **not participate in group work** are slightly more common.
- This could mean that **group work is optional** or that **many students prefer to study individually**.
- Important later to check if group work participation impacts results.

2.2 2. absences Distribution

- <1% absence dominates, followed by <25%, <50%, and >50%.
- Most students **have very low absence rates** (good attendance).
- Attendance seems generally **good** across the dataset.
- Important: students with high absences (>50%) might have lower final results — to verify later.

2.3 3. `schedule` Distribution

- “matin” (morning classes) are much more common than “soir” (evening classes).
- Majority of the courses happen **in the morning**.
- Maybe evening students are special cases (working students?), could affect performance — needs checking.

2.4 4. `class_participation` Distribution

- “un petit peu” (a little participation) is the most common, followed by “moyen” (average participation), and fewer “étudiant idéal” (ideal students).
- Most students **only participate a little** in class.
- Very **few are model students** (“étudiant idéal”).
- Important: more participation might correlate with better results — to check later.

2.5 5. `exam_content` Distribution

- “fait parti du cours enseigner par le prof” dominates heavily, “mélange” and “exam de la nasa” are less frequent.
- Most exams **are directly based on course material taught**.
- Some exams (“mélange” or “exam de la nasa”) may be harder or unexpected — worth checking if those students performed differently.


2.6 6. `result` (Target) Distribution

- “valide” is extremely dominant, while “valide après rattrapage”, “admis par compensation” are fewer, and “non valide” and “admis par dérogation” are very rare.
- **Most students pass** their courses normally.
- Very few fail completely (**non valide**) or pass exceptionally (**dérogation**).
- The dataset is **imbalanced** toward positive results → class imbalance techniques might be necessary during modeling.

```
[202]: # Select only numerical columns
numerical_cols = ['explanation_quality', 'final_grade', 'study_hours']

# Compute the correlation matrix
corr = data[numerical_cols].corr()

# Plot the heatmap
plt.figure(figsize=(8,6))
sns.heatmap(corr, annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```




```
[203]: # List of numerical columns
numerical_cols = ['explanation_quality', 'final_grade', 'study_hours']

# Create boxplots for each numerical feature against 'result'
plt.figure(figsize=(20,6))

for idx, col in enumerate(numerical_cols):
    plt.subplot(1, 3, idx + 1) # 1 row, 3 columns
    sns.boxplot(x='result', y=col, data=data)
    plt.title(f"{col} vs Result")
    plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

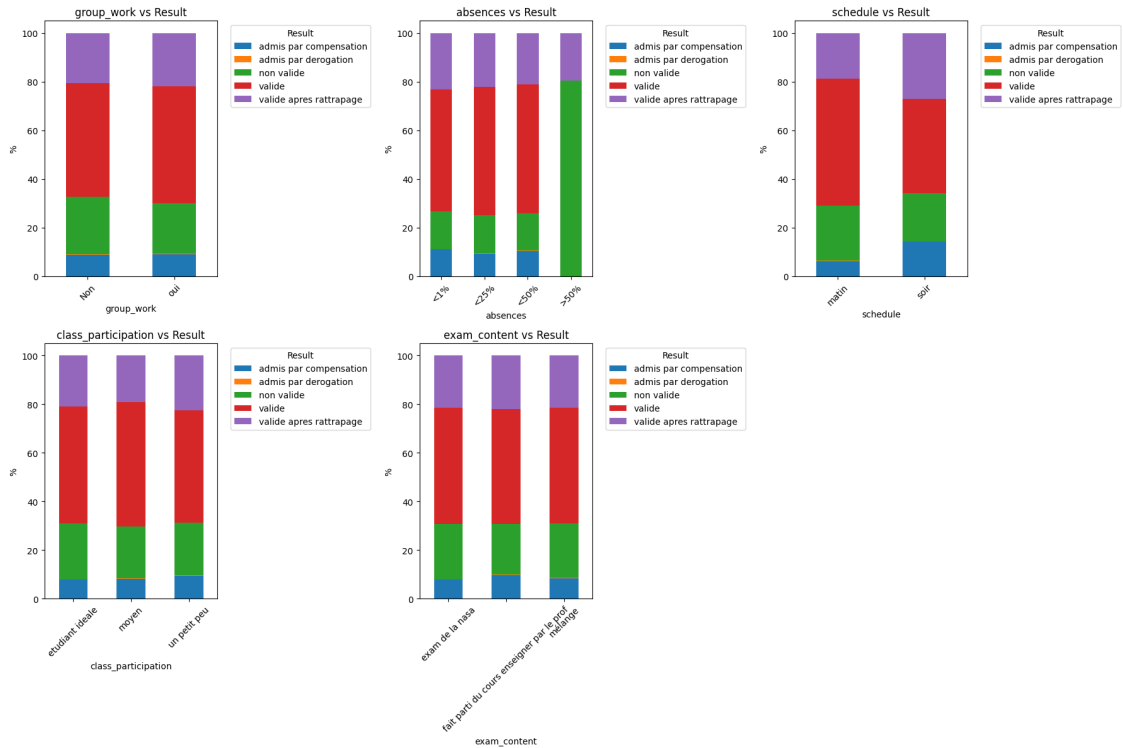


```
[204]: # List of categorical features
categorical_cols = ['group_work', 'absences', 'schedule', 'class_participation', 'exam_content']

# Create subplots
plt.figure(figsize=(18, 12))


for idx, col in enumerate(categorical_cols):
    plt.subplot(2, 3, idx + 1) # 2 rows, 3 columns
    crosstab = pd.crosstab(data[col], data['result'], normalize='index') * 100
    crosstab.plot(kind='bar', stacked=True, ax=plt.gca())
    plt.title(f"{col} vs Result")
    plt.ylabel('%')
    plt.xticks(rotation=45)
    plt.legend(title='Result', bbox_to_anchor=(1.05, 1), loc='upper left')

plt.tight_layout()
plt.show()
```



```
[205]: # Countplot of result
plt.figure(figsize=(8,6))
sns.countplot(x='result', data=data, order=data['result'].value_counts().index)
plt.title('Target Class Distribution (Result)')
plt.xlabel('Result')
plt.ylabel('Number of Students')
plt.xticks(rotation=45)
plt.show()

# Print percentage distribution
print(data['result'].value_counts(normalize=True) * 100)
```

```
result
valide          47.52
non valide      21.68
valide apres rattrapage  21.66
admis par compensation    8.96
admis par derogation     0.18
Name: proportion, dtype: float64
```

[]:

3 cleaning the study hours column

```
[206]: def clean_study_hours(value):
        if pd.isna(value):
            return None
        # Extract the first number found
        match = re.search(r'\d+', str(value))
        if match:
            return int(match.group())
        else:
            return None
```

```
[207]: data['study_hours'] = data['study_hours'].apply(clean_study_hours)
```

```
[208]: data.head(5000)
```

```
[208]:
```

	explanation_quality	final_grade	study_hours	absences	group_work	\
0	5	9.46	4	<25%	oui	
1	9	5.13	21	<25%	Non	
2	7	9.67	18	<50%	Non	
3	6	13.24	0	<25%	oui	
4	3	12.28	48	<50%	oui	
...	
4995	8	16.37	30	<50%	Non	
4996	3	12.55	34	<50%	oui	
4997	5	10.69	30	>50%	oui	
4998	7	14.72	26	<25%	Non	
4999	6	15.33	22	<25%	oui	

	schedule	class_participation	exam_content	\
0	matin	un petit peu	mélange	
1	matin	moyen	fait parti du cours enseigner par le prof	
2	matin	un petit peu	fait parti du cours enseigner par le prof	
3	matin	un petit peu	mélange	
4	matin	un petit peu	fait parti du cours enseigner par le prof	
...	
4995	soir	moyen	fait parti du cours enseigner par le prof	
4996	matin	un petit peu	mélange	
4997	matin	un petit peu	exam de la nasa	
4998	soir	un petit peu	mélange	
4999	matin	un petit peu	exam de la nasa	

	result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage

```

4             valide
...
4995  admis par compensation
4996  valide apres rattrapage
4997             non valide
4998             valide
4999             valide

```

[5000 rows x 9 columns]

```

[209]: for column in data.columns:
        unique_values = data[column].unique()
        print(f"Column: {column}")
        print(f"Unique values ({len(unique_values)}): {unique_values}\n")

```

Column: explanation_quality

Unique values (10): [5 9 7 6 3 2 8 1 4 10]

Column: final_grade

Unique values (1106): [9.46 5.13 9.67 ... 15.4 5.79 6.48]

Column: study_hours

Unique values (50): [4 21 18 0 48 29 44 20 27 37 2 26 15 30 35 25 23 32 17 22
31 3 34 24
28 16 40 42 19 38 12 36 13 33 14 9 7 6 11 8 41 39 5 1 10 47 43 45
50 46]

Column: absences

Unique values (4): ['<25%' '<50%' '<1%' '>50%']

Column: group_work

Unique values (2): ['oui' 'Non']

Column: schedule

Unique values (2): ['matin' 'soir']

Column: class_participation

Unique values (3): ['un petit peu' 'moyen' 'etudiant ideale']

Column: exam_content

Unique values (3): ['mélange' 'fait parti du cours enseigner par le prof' 'exam de la nasa']

Column: result

Unique values (5): ['non valide' 'valide' 'valide apres rattrapage' 'admis par compensation'
'admis par derogation']

```
[210]: rename_mapping = {
        'explanation_quality': 'Study Quality',
        'final_grade': 'Note',
        'result': 'Result',
        'study_hours': 'Autoformation',
        'absences': 'Absence',
        'group_work': 'Group Study',
        'exam_content': 'Alignement with Lecture',
        'schedule': 'Horaire',
        'class_participation': 'Class Engagement'
    }

    # Rename columns
    structured_data = data.rename(columns=rename_mapping)
```

```
[211]: structured_data.head()
```

```
[211]:
```

	Study Quality	Note	Autoformation	Absence	Group Study	Horaire	\
0	5	9.46	4	<25%	oui	matin	
1	9	5.13	21	<25%	Non	matin	
2	7	9.67	18	<50%	Non	matin	
3	6	13.24	0	<25%	oui	matin	
4	3	12.28	48	<50%	oui	matin	

	Class Engagement	Alignement with Lecture	\
0	un petit peu		mélange
1	moyen	fait parti du cours enseigner par le prof	
2	un petit peu	fait parti du cours enseigner par le prof	
3	un petit peu		mélange
4	un petit peu	fait parti du cours enseigner par le prof	

	Result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage
4	valide

4 plots


```
[212]: # Select numerical columns
numerical_cols = ['explanation_quality', 'final_grade', 'study_hours']

#cast final grade column to float
data['final_grade'] = data['final_grade'].astype(float)
```

```
# Create subplots
plt.figure(figsize=(15, 5))

for idx, col in enumerate(numerical_cols):
    plt.subplot(1, 3, idx + 1) # 1 row, 3 columns, plot number idx+1
    plt.hist(data[col], bins=30, edgecolor='black')
    plt.title(f"{col} Distribution")
    plt.xlabel(col)
    plt.ylabel("Frequency")

plt.tight_layout()
plt.show()
```



5 feature engineering

why binary encoding ?

```
[213]: # 1. Encode group_work (yes/no --> 1/0)
def encode_group_work(df: pd.DataFrame) -> pd.DataFrame:
    df['Group Study'] = df['Group Study'].map({'oui': 1, 'Non': 0})
    return df
```

```
[214]: encode_group_work(structured_data)
```

```
[214]:
```

	Study Quality	Note	Autoformation	Absence	Group Study	Horaire	\
0	5	9.46	4	<25%	1	matin	
1	9	5.13	21	<25%	0	matin	
2	7	9.67	18	<50%	0	matin	
3	6	13.24	0	<25%	1	matin	
4	3	12.28	48	<50%	1	matin	
...		
4995	8	16.37	30	<50%	0	soir	
4996	3	12.55	34	<50%	1	matin	

4997	5	10.69	30	>50%	1	matin
4998	7	14.72	26	<25%	0	soir
4999	6	15.33	22	<25%	1	matin

	Class Engagement	Alignement with Lecture	\
0	un petit peu		mélange
1	moyen	fait parti du cours enseigner par le prof	
2	un petit peu	fait parti du cours enseigner par le prof	
3	un petit peu		mélange
4	un petit peu	fait parti du cours enseigner par le prof	
...
4995	moyen	fait parti du cours enseigner par le prof	
4996	un petit peu		mélange
4997	un petit peu	exam de la nasa	
4998	un petit peu		mélange
4999	un petit peu	exam de la nasa	

	Result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage
4	valide
...	...
4995	admis par compensation
4996	valide apres rattrapage
4997	non valide
4998	valide
4999	valide

[5000 rows x 9 columns]

why ordinal encoding

```
[215]: # 2. Encode absences (<1%, <25%, etc.) as ordered categories
def encode_absences(df: pd.DataFrame) -> pd.DataFrame:
    absence_order = ['<1%', '<25%', '<50%', '>50%']
    ord_encoder = OrdinalEncoder(categories=[absence_order])
    df['Absence'] = ord_encoder.fit_transform(df[['Absence']])
    return df
```

```
[216]: encode_absences(structured_data)
```

	Study Quality	Note	Autoformation	Absence	Group Study	Horaires	\
0	5	9.46	4	1.0	1	matin	
1	9	5.13	21	1.0	0	matin	
2	7	9.67	18	2.0	0	matin	
3	6	13.24	0	1.0	1	matin	

4		3	12.28		48	2.0		1	matin
...		
4995		8	16.37		30	2.0		0	soir
4996		3	12.55		34	2.0		1	matin
4997		5	10.69		30	3.0		1	matin
4998		7	14.72		26	1.0		0	soir
4999		6	15.33		22	1.0		1	matin

	Class Engagement	Alignement with Lecture	\
0	un petit peu		mélange
1	moyen	fait parti du cours enseigner par le prof	
2	un petit peu	fait parti du cours enseigner par le prof	
3	un petit peu		mélange
4	un petit peu	fait parti du cours enseigner par le prof	
...
4995	moyen	fait parti du cours enseigner par le prof	
4996	un petit peu		mélange
4997	un petit peu	exam de la nasa	
4998	un petit peu		mélange
4999	un petit peu	exam de la nasa	

	Result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage
4	valide
...	...
4995	admis par compensation
4996	valide apres rattrapage
4997	non valide
4998	valide
4999	valide

[5000 rows x 9 columns]

why binary encoding

```
[217]: # 3. Encode schedule (matin/soir) as binary 0/1
def encode_schedule(df: pd.DataFrame) -> pd.DataFrame:
    df['Horaire'] = df['Horaire'].map({'matin': 0, 'soir': 1})
    return df
```

```
[218]: encode_schedule(structured_data)
```

	Study Quality	Note	Autoformation	Absence	Group Study	Horaire	\
0	5	9.46	4	1.0	1	0	
1	9	5.13	21	1.0	0	0	

2	7	9.67	18	2.0	0	0
3	6	13.24	0	1.0	1	0
4	3	12.28	48	2.0	1	0
...
4995	8	16.37	30	2.0	0	1
4996	3	12.55	34	2.0	1	0
4997	5	10.69	30	3.0	1	0
4998	7	14.72	26	1.0	0	1
4999	6	15.33	22	1.0	1	0

	Class Engagement	Alignement with Lecture \
0	un petit peu	mélange
1	moyen	fait parti du cours enseigner par le prof
2	un petit peu	fait parti du cours enseigner par le prof
3	un petit peu	mélange
4	un petit peu	fait parti du cours enseigner par le prof
...
4995	moyen	fait parti du cours enseigner par le prof
4996	un petit peu	mélange
4997	un petit peu	exam de la nasa
4998	un petit peu	mélange
4999	un petit peu	exam de la nasa

	Result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage
4	valide
...	...
4995	admis par compensation
4996	valide apres rattrapage
4997	non valide
4998	valide
4999	valide

[5000 rows x 9 columns]

why ordinal encoding

```
[219]: # 4. Encode class participation (un petit peu < moyen < etudiant ideale)
def encode_class_participation(df: pd.DataFrame) -> pd.DataFrame:
    participation_order = ['un petit peu', 'moyen', 'etudiant ideale']
    ord_encoder = OrdinalEncoder(categories=[participation_order])
    df['Class Engagement'] = ord_encoder.fit_transform(df[['Class Engagement']])
    return df
```

```
[220]: encode_class_participation(structured_data)
```



```
[220]:
```

	Study Quality	Note	Autoformation	Absence	Group Study	Horaire	\
0	5	9.46	4	1.0	1	0	
1	9	5.13	21	1.0	0	0	
2	7	9.67	18	2.0	0	0	
3	6	13.24	0	1.0	1	0	
4	3	12.28	48	2.0	1	0	
...	
4995	8	16.37	30	2.0	0	1	
4996	3	12.55	34	2.0	1	0	
4997	5	10.69	30	3.0	1	0	
4998	7	14.72	26	1.0	0	1	
4999	6	15.33	22	1.0	1	0	

	Class Engagement	Alignement with Lecture	\
0	0.0	mélange	
1	1.0	fait parti du cours enseigner par le prof	
2	0.0	fait parti du cours enseigner par le prof	
3	0.0	mélange	
4	0.0	fait parti du cours enseigner par le prof	
...	
4995	1.0	fait parti du cours enseigner par le prof	
4996	0.0	mélange	
4997	0.0	exam de la nasa	
4998	0.0	mélange	
4999	0.0	exam de la nasa	

	Result
0	non valide
1	non valide
2	valide
3	valide apres rattrapage
4	valide
...	...
4995	admis par compensation
4996	valide apres rattrapage
4997	non valide
4998	valide
4999	valide

[5000 rows x 9 columns]

why ordinal encoding

```
[221]: def onehot_encode_exam_content(df: pd.DataFrame) -> pd.DataFrame:
        # Ensure clean text (remove extra spaces)
        if df['Alignement with Lecture'].dtype == 'object':
            df['Alignement with Lecture'] = df['Alignement with Lecture'].str.
            ↪strip()
```

```

# Perform One-Hot Encoding
onehot = pd.get_dummies(
    df['Alignement with Lecture'],
    prefix='Alignement',
    dtype=int, # ensures 0/1 integers not booleans,
    drop_first=True # to avoid dummy variable trap
)

# Drop original and merge one-hot columns
df = df.drop(columns=['Alignement with Lecture'])
df = pd.concat([df, onehot], axis=1)

return df

```

```
[222]: structured_data = onehot_encode_exam_content(structured_data)
```

```
[223]: structured_data.head(5000)
```

```

[223]:
      Study Quality    Note  Autoformation  Absence  Group Study  Horaire  \
0           5    9.46           4         1.0         1         0
1           9    5.13          21         1.0         0         0
2           7    9.67          18         2.0         0         0
3           6   13.24           0         1.0         1         0
4           3   12.28          48         2.0         1         0
...         ...    ...         ...         ...         ...         ...
4995         8   16.37          30         2.0         0         1
4996         3   12.55          34         2.0         1         0
4997         5   10.69          30         3.0         1         0
4998         7   14.72          26         1.0         0         1
4999         6   15.33          22         1.0         1         0

      Class Engagement      Result  \
0           0.0      non valide
1           1.0      non valide
2           0.0       valide
3           0.0  valide apres rattrapage
4           0.0       valide
...         ...         ...
4995         1.0  admis par compensation
4996         0.0  valide apres rattrapage
4997         0.0      non valide
4998         0.0       valide
4999         0.0       valide

      Alignement_fait parti du cours enseigner par le prof  Alignement_mélange
0                                                         0                 1

```

1		1	0
2		1	0
3		0	1
4		1	0
...
4995		1	0
4996		0	1
4997		0	0
4998		0	1
4999		0	0

[5000 rows x 10 columns]

```
[224]: def encode_target(df: pd.DataFrame) -> pd.DataFrame:
        label_encoder = LabelEncoder()
        df['Result'] = label_encoder.fit_transform(df['Result'])
        return df
```

```
[225]: encode_target(structured_data)
```

```
[225]:
```

	Study Quality	Note	Autoformation	Absence	Group Study	Horaire	\
0	5	9.46	4	1.0	1	0	
1	9	5.13	21	1.0	0	0	
2	7	9.67	18	2.0	0	0	
3	6	13.24	0	1.0	1	0	
4	3	12.28	48	2.0	1	0	
...	
4995	8	16.37	30	2.0	0	1	
4996	3	12.55	34	2.0	1	0	
4997	5	10.69	30	3.0	1	0	
4998	7	14.72	26	1.0	0	1	
4999	6	15.33	22	1.0	1	0	

	Class Engagement	Result	\
0	0.0	2	
1	1.0	2	
2	0.0	3	
3	0.0	4	
4	0.0	3	
...	
4995	1.0	0	
4996	0.0	4	
4997	0.0	2	
4998	0.0	3	
4999	0.0	3	

Alignement_fait parti du cours enseigner par le prof Alignement_mélange

0	0	1
1	1	0
2	1	0
3	0	1
4	1	0
...
4995	1	0
4996	0	1
4997	0	0
4998	0	1
4999	0	0

[5000 rows x 10 columns]

6 Separate features and target

```
[226]: X = structured_data.drop(columns=['Result']) # Features (everything except the
        ↪target)
        y = structured_data['Result'] # Target variable
```

```
[227]: print(y)
```

```
0      2
1      2
2      3
3      4
4      3
..
4995   0
4996   4
4997   2
4998   3
4999   3
```

Name: Result, Length: 5000, dtype: int64

7 Split X and y into training and testing sets

```
[228]: X_train, X_test, y_train, y_test = train_test_split(
        X, y,
        test_size=0.3,      # 30% test, 70% train
        random_state=42,    # random state for reproducibility
        stratify=y          # keep the same distribution of result classes
    )

    # Let's check the shapes
    print("Training set:", X_train.shape, y_train.shape)
```

```
print("Testing set:", X_test.shape, y_test.shape)
```

Training set: (3500, 9) (3500,)

Testing set: (1500, 9) (1500,)

```
[229]: try:
        results
    except NameError:
        results = []

    print("="*50)
    print("Training: Logistic Regression")
    print("="*50)

    start_train = time.time()
    log_reg = LogisticRegression(max_iter=1000)
    param_grid_log_reg = {
        'C': [0.01, 0.1, 1, 10],
        'solver': ['lbfgs', 'liblinear']
    }
    grid_log = GridSearchCV(log_reg, param_grid_log_reg, cv=3, n_jobs=-1,
        ↪scoring='accuracy')
    grid_log.fit(X_train, y_train)
    best_log = grid_log.best_estimator_
    train_duration = round(time.time() - start_train, 2)

    start_pred = time.time()
    y_pred_log = best_log.predict(X_test)
    pred_duration = round(time.time() - start_pred, 2)

    acc_log = accuracy_score(y_test, y_pred_log)
    print(f"Best Params: {grid_log.best_params_}")
    print(f"Accuracy: {acc_log:.4f}")
    print(classification_report(y_test, y_pred_log, zero_division=0))

    results.append({
        'Model': 'Logistic Regression',
        'Best Params': grid_log.best_params_,
        'Accuracy': acc_log,
        'Training Time (s)': train_duration,
        'Prediction Time (s)': pred_duration
    })
```

```
=====
Training: Logistic Regression
=====
```

```
Best Params: {'C': 0.01, 'solver': 'lbfgs'}
```

```
Accuracy: 0.5727
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.59	0.62	0.61	325
3	0.57	0.92	0.70	713
4	0.00	0.00	0.00	325
accuracy			0.57	1500
macro avg	0.23	0.31	0.26	1500
weighted avg	0.40	0.57	0.47	1500

c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[230]: print("="*50)
print("Training: KNN")
print("="*50)

start_train = time.time()
knn = KNeighborsClassifier()
param_grid_knn = {
    'n_neighbors': [3, 5, 7],
    'weights': ['uniform', 'distance']
}
grid_knn = GridSearchCV(knn, param_grid_knn, cv=3, n_jobs=-1,
    ↪scoring='accuracy')
grid_knn.fit(X_train, y_train)
best_knn = grid_knn.best_estimator_
train_duration = round(time.time() - start_train, 2)

start_pred = time.time()
y_pred_knn = best_knn.predict(X_test)
pred_duration = round(time.time() - start_pred, 2)

acc_knn = accuracy_score(y_test, y_pred_knn)
print(f"Best Params: {grid_knn.best_params_}")
print(f"Accuracy: {acc_knn:.4f}")
```

```

print(classification_report(y_test, y_pred_knn, zero_division=0))

results.append({
    'Model': 'KNN',
    'Best Params': grid_knn.best_params_,
    'Accuracy': acc_knn,
    'Training Time (s)': train_duration,
    'Prediction Time (s)': pred_duration
})

```

```

=====
Training: KNN
=====
Best Params: {'n_neighbors': 7, 'weights': 'uniform'}
Accuracy: 0.6093

```

	precision	recall	f1-score	support
0	0.07	0.02	0.03	134
1	0.00	0.00	0.00	3
2	0.76	0.77	0.76	325
3	0.62	0.88	0.73	713
4	0.29	0.10	0.15	325
accuracy			0.61	1500
macro avg	0.35	0.35	0.33	1500
weighted avg	0.53	0.61	0.55	1500

```

[231]: print("="*50)
print("Training: SVM")
print("="*50)

start_train = time.time()
svm = SVC()
param_grid_svm = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'rbf', 'poly'], # Added 'poly' for polynomial kernel
    'degree': [2, 3, 4] # Degree of the polynomial kernel
}
grid_svm = GridSearchCV(svm, param_grid_svm, cv=3, n_jobs=-1,
    ↪scoring='accuracy')
grid_svm.fit(X_train, y_train)
best_svm = grid_svm.best_estimator_
train_duration = round(time.time() - start_train, 2)

start_pred = time.time()
y_pred_svm = best_svm.predict(X_test)

```

```

pred_duration = round(time.time() - start_pred, 2)

acc_svm = accuracy_score(y_test, y_pred_svm)
print(f"Best Params: {grid_svm.best_params_}")
print(f"Accuracy: {acc_svm:.4f}")
print(classification_report(y_test, y_pred_svm, zero_division=0))

results.append({
    'Model': 'SVM',
    'Best Params': grid_svm.best_params_,
    'Accuracy': acc_svm,
    'Training Time (s)': train_duration,
    'Prediction Time (s)': pred_duration
})

```

=====
Training: SVM
=====

Best Params: {'C': 10, 'degree': 2, 'kernel': 'rbf'}

Accuracy: 0.6320

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.68	0.83	0.75	325
3	0.62	0.95	0.75	713
4	0.00	0.00	0.00	325
accuracy			0.63	1500
macro avg	0.26	0.36	0.30	1500
weighted avg	0.44	0.63	0.52	1500

```

[232]: print("="*50)
print("Training: Decision Tree")
print("="*50)

# 1. Training + Grid Search
start_train = time.time()
tree = DecisionTreeClassifier()
param_grid_tree = {
    'max_depth': [1,2,3,4,5, 10, 20,50],
    'criterion': ['gini', 'entropy']
}
grid_tree = GridSearchCV(tree, param_grid_tree, cv=3, n_jobs=-1,
    ↪scoring='accuracy')
grid_tree.fit(X_train, y_train)

```



```

best_tree = grid_tree.best_estimator_
train_duration = round(time.time() - start_train, 2)

# 2. Prediction
start_pred = time.time()
y_pred_tree = best_tree.predict(X_test)
pred_duration = round(time.time() - start_pred, 2)

# 3. Evaluation
acc_tree = accuracy_score(y_test, y_pred_tree)
print(f"Best Params: {grid_tree.best_params_}")
print(f"Accuracy: {acc_tree:.4f}")
print(classification_report(y_test, y_pred_tree, zero_division=0))

# 4. Save results
results.append({
    'Model': 'Decision Tree',
    'Best Params': grid_tree.best_params_,
    'Accuracy': acc_tree,
    'Training Time (s)': train_duration,
    'Prediction Time (s)': pred_duration
})

```

```

=====
Training: Decision Tree
=====
Best Params: {'criterion': 'gini', 'max_depth': 3}
Accuracy: 0.6853

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.80	0.97	0.88	325
3	0.65	1.00	0.78	713
4	0.00	0.00	0.00	325
accuracy			0.69	1500
macro avg	0.29	0.39	0.33	1500
weighted avg	0.48	0.69	0.56	1500

```

[233]: print("="*50)
print("Training: Random Forest")
print("="*50)

# 1. Training + Grid Search
start_train = time.time()

```

```

rf = RandomForestClassifier()
param_grid_rf = {
    'n_estimators': [50,100, 200,300],
    'max_depth': [1,2,3,4,5,20,50,100],
}
grid_rf = GridSearchCV(rf, param_grid_rf, cv=3, n_jobs=-1, scoring='accuracy')
grid_rf.fit(X_train, y_train)
best_rf = grid_rf.best_estimator_
train_duration = round(time.time() - start_train, 2)

# 2. Prediction
start_pred = time.time()
y_pred_rf = best_rf.predict(X_test)
pred_duration = round(time.time() - start_pred, 2)

# 3. Evaluation
acc_rf = accuracy_score(y_test, y_pred_rf)
print(f"Best Params: {grid_rf.best_params_}")
print(f"Accuracy: {acc_rf:.4f}")
print(classification_report(y_test, y_pred_rf, zero_division=0))

# 4. Save results
results.append({
    'Model': 'Random Forest',
    'Best Params': grid_rf.best_params_,
    'Accuracy': acc_rf,
    'Training Time (s)': train_duration,
    'Prediction Time (s)': pred_duration
})

```

```

=====
Training: Random Forest
=====
Best Params: {'max_depth': 5, 'n_estimators': 50}
Accuracy: 0.6853

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.80	0.97	0.88	325
3	0.65	1.00	0.78	713
4	0.00	0.00	0.00	325
accuracy			0.69	1500
macro avg	0.29	0.39	0.33	1500
weighted avg	0.48	0.69	0.56	1500

```
[234]: print("="*50)
print("Training: LightGBM")
print("="*50)

# 1. Training + Grid Search
start_train = time.time()
lgbm = LGBMClassifier()
param_grid_lgbm = {
    'n_estimators': [50,100, 200,300],
    'learning_rate': [0.05, 0.1]
}
grid_lgbm = GridSearchCV(lgbm, param_grid_lgbm, cv=3, n_jobs=-1,
    scoring='accuracy')
grid_lgbm.fit(X_train, y_train)
best_lgbm = grid_lgbm.best_estimator_
train_duration = round(time.time() - start_train, 2)

# 2. Prediction
start_pred = time.time()
y_pred_lgbm = best_lgbm.predict(X_test)
pred_duration = round(time.time() - start_pred, 2)

# 3. Evaluation
acc_lgbm = accuracy_score(y_test, y_pred_lgbm)
print(f"Best Params: {grid_lgbm.best_params_}")
print(f"Accuracy: {acc_lgbm:.4f}")
print(classification_report(y_test, y_pred_lgbm, zero_division=0))

# 4. Save results
results.append({
    'Model': 'LightGBM',
    'Best Params': grid_lgbm.best_params_,
    'Accuracy': acc_lgbm,
    'Training Time (s)': train_duration,
    'Prediction Time (s)': pred_duration
})
```

```
=====
Training: LightGBM
=====
```

```
c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\utils\_tags.py:354: FutureWarning: The LGBMClassifier or
classes from which it inherits use `_get_tags` and `_more_tags`. Please define
the `__sklearn_tags__` method, or inherit from `sklearn.base.BaseEstimator`
and/or other appropriate mixins such as `sklearn.base.TransformerMixin`,
`sklearn.base.ClassifierMixin`, `sklearn.base.RegressorMixin`, and
`sklearn.base.OutlierMixin`. From scikit-learn 1.7, not defining
```

```
`__sklearn_tags__` will raise an error.  
warnings.warn(  

```

```
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines  
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of  
testing was 0.000243 seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 328
```

```
[LightGBM] [Info] Number of data points in the train set: 3500, number of used  
features: 9
```

```
[LightGBM] [Info] Start training from score -2.411125
```

```
[LightGBM] [Info] Start training from score -6.368759
```

```
[LightGBM] [Info] Start training from score -1.528516
```

```
[LightGBM] [Info] Start training from score -0.744140
```

```
[LightGBM] [Info] Start training from score -1.529835
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

```
Best Params: {'learning_rate': 0.05, 'n_estimators': 50}
```

```
Accuracy: 0.6840
```

	precision	recall	f1-score	support
0	0.25	0.01	0.01	134
1	0.00	0.00	0.00	3
2	0.80	0.97	0.88	325
3	0.65	0.98	0.78	713
4	0.35	0.03	0.05	325
accuracy			0.68	1500
macro avg	0.41	0.40	0.35	1500
weighted avg	0.58	0.68	0.57	1500

```
[235]: # Create a DataFrame from the collected results  
results_df = pd.DataFrame(results)  
  
# Sort the models by their Accuracy (optional)  
results_df = results_df.sort_values(by='Accuracy', ascending=False)  
  
# Display the table  
print("\n\nFinal Model Performance Summary:\n")  
print(results_df.to_string(index=False))
```

Final Model Performance Summary:

Model	Best Params	Accuracy
Training Time (s)	Prediction Time (s)	
Random Forest	{ 'max_depth': 3, 'n_estimators': 50 }	0.686000

26.51		0.01		
	Random Forest		{'max_depth': 5, 'n_estimators': 50}	0.685333
25.28		0.01		
	Random Forest		{'max_depth': 5, 'n_estimators': 50}	0.685333
20.39		0.02		
	Decision Tree		{'criterion': 'gini', 'max_depth': 3}	0.685333
0.47		0.00		
	Decision Tree		{'criterion': 'gini', 'max_depth': 3}	0.685333
0.45		0.00		
	Random Forest		{'max_depth': 10, 'n_estimators': 200}	0.684000
6.24		0.05		
	LightGBM		{'learning_rate': 0.05, 'n_estimators': 50}	0.684000
14.71		0.02		
	LightGBM		{'learning_rate': 0.05, 'n_estimators': 50}	0.684000
22.15		0.02		
	Random Forest		{'max_depth': 10, 'n_estimators': 300}	0.683333
18.33		0.11		
	Random Forest		{'max_depth': 10, 'n_estimators': 200}	0.683333
5.72		0.05		
	Random Forest		{'max_depth': 10, 'n_estimators': 200}	0.683333
5.65		0.05		
	Decision Tree		{'criterion': 'entropy', 'max_depth': 5}	0.682000
0.29		0.00		
	Decision Tree		{'criterion': 'entropy', 'max_depth': 5}	0.682000
0.27		0.00		
	Decision Tree		{'criterion': 'entropy', 'max_depth': 5}	0.682000
0.30		0.00		
	Decision Tree		{'criterion': 'entropy', 'max_depth': 5}	0.682000
0.26		0.00		
	LightGBM		{'learning_rate': 0.05, 'n_estimators': 100}	0.668000
17.00		0.04		
	LightGBM		{'learning_rate': 0.05, 'n_estimators': 100}	0.668000
10.76		0.04		
	LightGBM		{'learning_rate': 0.05, 'n_estimators': 100}	0.668000
28.19		0.04		
	Random Forest		{'max_depth': 20, 'n_estimators': 300}	0.663333
21.81		0.15		
	SVM		{'C': 10, 'kernel': 'rbf'}	0.632000
201.91		0.52		
	SVM		{'C': 10, 'kernel': 'rbf'}	0.632000
196.73		0.56		
	SVM		{'C': 10, 'degree': 2, 'kernel': 'rbf'}	0.632000
386.40		0.71		
	SVM		{'C': 10, 'kernel': 'rbf'}	0.632000
140.96		0.43		
	SVM		{'C': 10, 'kernel': 'rbf'}	0.632000
142.60		0.48		
	SVM		{'C': 10, 'degree': 2, 'kernel': 'rbf'}	0.632000

394.87	0.50		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
0.68	0.17		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
1.08	0.14		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
1.50	0.27		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
0.69	0.56		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
1.00	0.15		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
1.27	0.16		
	KNN	{'n_neighbors': 7, 'weights': 'uniform'}	0.609333
0.71	0.08		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
14.34	0.00		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
5.38	0.00		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
9.41	0.00		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
9.76	0.00		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
6.75	0.00		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
12.32	0.01		
Logistic Regression		{'C': 0.01, 'solver': 'lbfgs'}	0.572667
7.71	0.01		

```
[238]: import time
import pandas as pd
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, classification_report

# Models
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from lightgbm import LGBMClassifier
from xgboost import XGBClassifier

# Initialize results
results = []
```

```

# Define models and parameters
models_and_params = {
    'Logistic Regression': (LogisticRegression(max_iter=1000), {
        'C': [0.01, 0.1, 1, 10, 50],
        'solver': ['lbfgs', 'liblinear']
    }),
    'KNN': (KNeighborsClassifier(), {
        'n_neighbors': list(range(3, 11)),
        'weights': ['uniform', 'distance']
    }),
    'SVM': (SVC(), {
        'C': [0.1, 1, 10, 100],
        'kernel': ['linear', 'rbf', 'poly'],
        'degree': [2, 3, 4]
    }),
    'Decision Tree': (DecisionTreeClassifier(), {
        'max_depth': [3, 5, 10, 20],
        'criterion': ['gini', 'entropy']
    }),
    'Random Forest': (RandomForestClassifier(), {
        'n_estimators': [50, 100, 200, 300],
        'max_depth': [5, 10, 20, 30]
    }),
    'LightGBM': (LGBMClassifier(), {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.05, 0.1]
    }),
}

# Train models
for model_name, (model, param_dist) in models_and_params.items():
    print("="*50)
    print(f"Training: {model_name}")
    print("="*50)

# 1. Train
start_train = time.time()
random_search = RandomizedSearchCV(
    model,
    param_distributions=param_dist,
    n_iter=20,          # Try 10 random combinations
    cv=3,
    n_jobs=-1,
    scoring='accuracy',
    random_state=42
)

```

```

)
random_search.fit(X_train, y_train)
best_model = random_search.best_estimator_
train_duration = round(time.time() - start_train, 2)

# 2. Predict
start_pred = time.time()
y_pred = best_model.predict(X_test)
pred_duration = round(time.time() - start_pred, 2)

# 3. Evaluate
acc = accuracy_score(y_test, y_pred)
print(f"Best Params: {random_search.best_params_}")
print(f"Accuracy: {acc:.4f}")
print(classification_report(y_test, y_pred, zero_division=0))

# 4. Save results
results.append({
    'Model': model_name,
    'Best Params': random_search.best_params_,
    'Accuracy': acc,
    'Training Time (s)': train_duration,
    'Prediction Time (s)': pred_duration
})

# Final Table
results_df = pd.DataFrame(results)
print("\n\nFinal Model Performance Summary:")
print(results_df.sort_values(by='Accuracy', ascending=False))

```

=====

Training: Logistic Regression

=====

c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\model_selection_search.py:317: UserWarning: The total space of parameters 10 is smaller than n_iter=20. Running 10 iterations. For exhaustive searches, use GridSearchCV.

warnings.warn(

c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-


```

regression
    n_iter_i = _check_optimize_result(
c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\model_selection\_search.py:317: UserWarning: The total space of
parameters 16 is smaller than n_iter=20. Running 16 iterations. For exhaustive
searches, use GridSearchCV.
    warnings.warn(

```

Best Params: {'solver': 'lbfgs', 'C': 0.01}
Accuracy: 0.5727

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.59	0.62	0.61	325
3	0.57	0.92	0.70	713
4	0.00	0.00	0.00	325
accuracy			0.57	1500
macro avg	0.23	0.31	0.26	1500
weighted avg	0.40	0.57	0.47	1500

=====
Training: KNN
=====

Best Params: {'weights': 'uniform', 'n_neighbors': 10}
Accuracy: 0.6113

	precision	recall	f1-score	support
0	0.03	0.01	0.01	134
1	0.00	0.00	0.00	3
2	0.78	0.74	0.76	325
3	0.60	0.91	0.73	713
4	0.28	0.07	0.11	325
accuracy			0.61	1500
macro avg	0.34	0.35	0.32	1500
weighted avg	0.52	0.61	0.54	1500

=====
Training: SVM
=====

Best Params: {'kernel': 'rbf', 'degree': 3, 'C': 100}
Accuracy: 0.6553

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3

2	0.72	0.90	0.80	325
3	0.63	0.97	0.77	713
4	0.00	0.00	0.00	325
accuracy			0.66	1500
macro avg	0.27	0.37	0.31	1500
weighted avg	0.46	0.66	0.54	1500

=====
Training: Decision Tree
=====

Best Params: {'max_depth': 3, 'criterion': 'gini'}
Accuracy: 0.6853

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.80	0.97	0.88	325
3	0.65	1.00	0.78	713
4	0.00	0.00	0.00	325
accuracy			0.69	1500
macro avg	0.29	0.39	0.33	1500
weighted avg	0.48	0.69	0.56	1500

=====
Training: Random Forest
=====

c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\model_selection_search.py:317: UserWarning: The total space of parameters 8 is smaller than n_iter=20. Running 8 iterations. For exhaustive searches, use GridSearchCV.

warnings.warn(

c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\model_selection_search.py:317: UserWarning: The total space of parameters 16 is smaller than n_iter=20. Running 16 iterations. For exhaustive searches, use GridSearchCV.

warnings.warn(

Best Params: {'n_estimators': 50, 'max_depth': 5}
Accuracy: 0.6853

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.80	0.97	0.88	325
3	0.65	1.00	0.78	713
4	0.00	0.00	0.00	325

accuracy			0.69	1500
macro avg	0.29	0.39	0.33	1500
weighted avg	0.48	0.69	0.56	1500

```
=====
Training: LightGBM
=====
```

```
c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\utils\_tags.py:354: FutureWarning: The LGBMClassifier or
classes from which it inherits use `_get_tags` and `_more_tags`. Please define
the `__sklearn_tags__` method, or inherit from `sklearn.base.BaseEstimator`
and/or other appropriate mixins such as `sklearn.base.TransformerMixin`,
`sklearn.base.ClassifierMixin`, `sklearn.base.RegressorMixin`, and
`sklearn.base.OutlierMixin`. From scikit-learn 1.7, not defining
`__sklearn_tags__` will raise an error.
```

```
warnings.warn(
```

```
c:\Users\admin\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\model_selection\_search.py:317: UserWarning: The total space of
parameters 9 is smaller than n_iter=20. Running 9 iterations. For exhaustive
searches, use GridSearchCV.
```

```
warnings.warn(
```

```
[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000217 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 328
```

```
[LightGBM] [Info] Number of data points in the train set: 3500, number of used
features: 9
```

```
[LightGBM] [Info] Start training from score -2.411125
```

```
[LightGBM] [Info] Start training from score -6.368759
```

```
[LightGBM] [Info] Start training from score -1.528516
```

```
[LightGBM] [Info] Start training from score -0.744140
```

```
[LightGBM] [Info] Start training from score -1.529835
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

```
Best Params: {'n_estimators': 50, 'learning_rate': 0.01}
```

```
Accuracy: 0.6853
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	134
1	0.00	0.00	0.00	3
2	0.80	0.97	0.88	325
3	0.64	1.00	0.78	713
4	0.00	0.00	0.00	325
accuracy			0.69	1500

macro avg	0.29	0.39	0.33	1500
weighted avg	0.48	0.69	0.56	1500

Final Model Performance Summary:

	Model	Best Params	Accuracy \
4	Random Forest	{'n_estimators': 50, 'max_depth': 5}	0.685333
3	Decision Tree	{'max_depth': 3, 'criterion': 'gini'}	0.685333
5	LightGBM	{'n_estimators': 50, 'learning_rate': 0.01}	0.685333
2	SVM	{'kernel': 'rbf', 'degree': 3, 'C': 100}	0.655333
1	KNN	{'weights': 'uniform', 'n_neighbors': 10}	0.611333
0	Logistic Regression	{'solver': 'lbfgs', 'C': 0.01}	0.572667

	Training Time (s)	Prediction Time (s)
4	16.75	0.01
3	0.19	0.00
5	13.22	0.04
2	440.98	0.68
1	2.44	0.26
0	4.60	0.00